

Synchronisation: Reader-Writer Problem

Master M1 MOSIG, Grenoble University

2014

I. The Reader-Writer Problem

We consider a situation in which a shared resource is concurrently accessed by two types of entities: the readers and the writers. The readers access the resource without modifying it. The writers, can alter the resource. To keep the resource in a consistent state, multiple readers may access it concurrently but writers should access it exclusively. In other terms, when a writer is using a resource, no reader or writer can access this resource.

The reader-writer problem is a classical synchronization problem. This access scheme is typically used when working with files or memory.

In the following we will consider the following functions:

- `init(reader_writer_t)` is the initialization function
- `begin_read(reader_writer_t)` and `end_read(reader_writer_t)` are the functions used by each reader before and after reading
- `begin_write(reader_writer_t)` et `end_write(reader_writer_t)` are the functions used by a writer before and after writing.

The structure `reader_writer_t` is to be defined. It must contain all the variables used by the synchronization primitives.

II. Implementing a thread-safe linked list

You are asked to implement a **thread-safe linked list**. A **non** thread-safe linked list implementation is included in the archive¹ (`linked_list.c`, `linked_list.h`).

In the linked list the `list_insert()` and `list_remove()` are considered to be *writings*, since they alter the linked list. `list_exists()` is considered to be *reading*.

The provided implementation includes correct calls to the `begin_read()`, `end_read()`, `begin_write()` and `end_write()` primitives. Your work consists in implementing these primitives so the linked list behaves correctly when accessed by multiple threads.

¹http://membres-lig.imag.fr/negrevergne/ens/assignment_two.tar.gz

```

Thread nmbr 0, won.
Thread nmbr 1, won.
THREAD 0 TIME: +0s 0.046ms, TYPE : BEGIN WRITE
THREAD 0 TIME: +0s 0.047ms, TYPE : END WRITE
THREAD 0 TIME: +0s 0.047ms, TYPE : BEGIN READ
THREAD 0 TIME: +0s 0.047ms, TYPE : END READ
THREAD 1 TIME: +0s 0.062ms, TYPE : BEGIN WRITE
THREAD 1 TIME: +0s 0.065ms, TYPE : END WRITE
THREAD 1 TIME: +0s 0.065ms, TYPE : BEGIN READ
THREAD 1 TIME: +0s 0.066ms, TYPE : END READ
THREAD 1 TIME: +0s 0.066ms, TYPE : BEGIN WRITE
THREAD 1 TIME: +0s 0.066ms, TYPE : END WRITE
THREAD 1 TIME: +0s 0.066ms, TYPE : BEGIN READ
THREAD 1 TIME: +0s 0.066ms, TYPE : END READ
THREAD 1 TIME: +0s 0.067ms, TYPE : BEGIN WRITE
THREAD 1 TIME: +0s 0.067ms, TYPE : END WRITE
THREAD 1 TIME: +0s 0.067ms, TYPE : BEGIN READ
THREAD 1 TIME: +0s 0.067ms, TYPE : END READ
THREAD 1 TIME: +0s 0.067ms, TYPE : BEGIN WRITE
THREAD 1 TIME: +0s 0.068ms, TYPE : END WRITE
THREAD 1 TIME: +0s 0.068ms, TYPE : BEGIN READ
THREAD 1 TIME: +0s 0.068ms, TYPE : END READ

```

Figure 1: output of : linked_list_simple_test

The archive also contains a simple test file `linked_list_simple_test.c`. Read it carefully. Any successful test must return `EXIT_SUCCESS`.

Question II.1: *Considering that no synchronization has been implemented yet, what result could you expect from this test. If you run the test, it won't actually fail, why ? What do you conclude on the validity of this test ?*

The code of `begin_read()`, `end_read()`, `begin_write()` and `end_write()` contains calls to the tracing module (`reader_writer_tracing.h`, `reader_writer_tracing.c`). This module allow one to record events and to prints them afterward.

Question II.2: *Fig. 1 is a possible output of `./linked_list_simple_test` 2. Using this output, how can you verify the consistency of the calls to `begin_read()`, `end_read()`, `begin_write()` and `end_write()`.*

Question II.3: *Write a naïve implementation of `begin_read()`, `end_read()`, `begin_write()` and `end_write()` in a new file named `reader_writer_1.c` using only mutexes².*

Question II.4: *Write a test program using the tracing module to check the consistency of the calls to `begin_read()`, `end_read()`, `begin_write()` and `end_write()`.*

²You must uncomment the corresponding line in the Makefile

Question II.5: *What could be the definition of efficiency in this particular problem ? In synchronization problems in general ? How would you evaluate the efficiency of this solution according to your previous definition ?*

III. Improving concurrency

The previous solution does not meet the requirements in terms of concurrency since two readers cannot read the list simultaneously.

Question III.1: *Implement new `begin_read()`, `end_read()`, `begin_write()` and `end_write()` functions in `reader_writer_2.c` to improve the concurrent capabilities of your thread-safe list. You may use conditions.*

Question III.2: *Write a test program that verify that readers can access the list concurrently. Note that this test will fail with your first implementation and pass with your second implementation.*

IV. Improving fairness

Question IV.1: *Introducing the events `WAITS TO READ` and `WAITS TO WRITE` for threads waiting in the functions `begin_read()` and `begin_write()`. Spot the fairness issue in Fig. 2.*

Question IV.2: *Discuss about the **fairness** of your previous solution.*

Question IV.3: *Build a fair solution for the reader writer problem in `reader_writer_3.c`. Write a corresponding test.
Hint : You must not mix the readers which arrived before a writer with the ones arriving after.*

```

THREAD 0 TIME: +0s 0.068ms, TYPE : BEGIN READ
THREAD 1 TIME: +0s 0.069ms, TYPE : WAITS TO WRITE
THREAD 2 TIME: +0s 0.070ms, TYPE : BEGIN READ
THREAD 0 TIME: +0s 0.071ms, TYPE : END READ

THREAD 0 TIME: +0s 0.083ms, TYPE : BEGIN READ
THREAD 2 TIME: +0s 0.085ms, TYPE : END READ

THREAD 2 TIME: +0s 0.123ms, TYPE : BEGIN READ
THREAD 0 TIME: +0s 0.127ms, TYPE : END READ

THREAD 0 TIME: +0s 0.143ms, TYPE : BEGIN READ
THREAD 2 TIME: +0s 0.145ms, TYPE : END READ

THREAD 2 TIME: +0s 0.161ms, TYPE : BEGIN READ
THREAD 0 TIME: +0s 0.168ms, TYPE : END READ
...
THREAD 2 TIME: +0s 0.183ms, TYPE : END READ
THREAD 1 TIME: +0s 0.184ms, TYPE : BEGIN WRITE
THREAD 1 TIME: +0s 0.187ms, TYPE : END WRITE

```

Figure 2: problematic behavior